deeplearning.ai

Comparing to human-level performance

Surpassing human-level performance

Surpassing human-level performance

Team of humans 0.5%

One human $~~~~\cancel{1\%}~~~~$ 0.1

Training error 0.6%

Dev error 0.8% 0.2

What is <u>avoidable bias?</u>

0.5% ← 0.1%
0.2%
1% 0.3%

0.3%

0.4%

Andrew Ng

Problems where ML significantly surpasses human-level performance

→ - Online advertising

→ - Product recommendations

→ - Logistics (predicting transit time)

→ - Loan approvals

Structual data

Not natural perception

Lots of data

{ - Speech recognition
  - Some image recognition
  - Medical
    - ECG, Skin cancer, ...

Andrew Ng